

Multiprocessor programming, DV2597/DV2606

Parallel Programming Models (Chapter 3)

Håkan Grahn

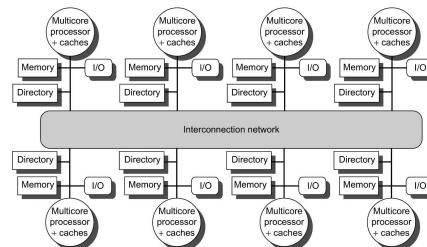
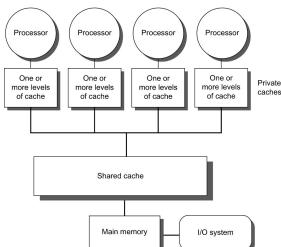
(some slides by G. Karypis, T. Rauber and G. Rünger)

Topic overview

- Models for parallel systems
- Parallelization of programs
- Levels of parallelism
- Data distribution for arrays
- Information exchange
- Parallel Matrix-Vector Multiplication

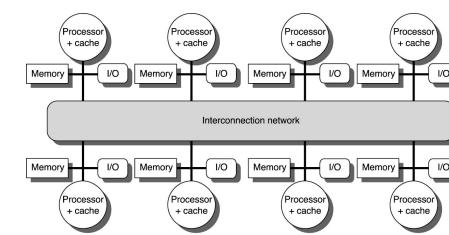
Shared-Address-Space Platforms

- The memory is accessible to all processors.
 - Processors interact by modifying data objects stored in this shared-address-space => Needs synchronization!
- If the access time of any memory word in the system is equal, the platform is classified as a uniform memory access (UMA), else, a non-uniform memory access (NUMA) machine.



Message-Passing Platforms

- These platforms comprise of a set of processors and their own (exclusive) memory.
 - Examples: Clustered workstations and non-shared-address-space multicomputers.
- These platforms are programmed using (variants of) send and receive primitives.
 - Libraries such as MPI and PVM provide such primitives.



Topic overview

- Models for parallel systems
- Parallelization of programs
- Levels of parallelism
- Data distribution for arrays
- Information exchange
- Parallel Matrix-Vector Multiplication

Models for parallel systems

Distinction according to level of abstraction

- **Parallel machine models:** lowest level of abstraction – hardware related description of the system
- **Parallel architectural models:** Abstraction of machine models – (topology, synchronous or asynchronous operation of the processor, SIMD or MIMD, memory organization)
- **Parallel computational models:** Extension of the architectural models, by which algorithms can be constructed and their costs can be considered, e.g., PRAM-Model (parallel random access machine)
- **Parallel programming models:** Description of a parallel system by describing the programming language and environment

Criteria for parallel programming models

- What kind of parallelism from the computation can be used?
(instruction level parallelism, function level, parallel loops)
- Has the programmer to specify the parallelism and how is the parallelism specified?
(explicit or implicit specification of parallelism)
- In which way has the programmer to specify the parallelism?
(e.g., independent tasks, managed by task pools or processes that are generated and have to communicate to each other)
- How is the execution of the parallel units organized?
(SIMD or SPMD, synchronous or asynchronous)
- How is the information exchange organized?
(communication with messages or by using shared variables)
- What kind of synchronization can be used?

Topic overview

- Models for parallel systems
- Parallelization of programs
- Levels of parallelism
- Data distribution for arrays
- Information exchange
- Parallel Matrix-Vector Multiplication

Parallelization of programs (I)

1. Decomposition (partitioning) of the computations

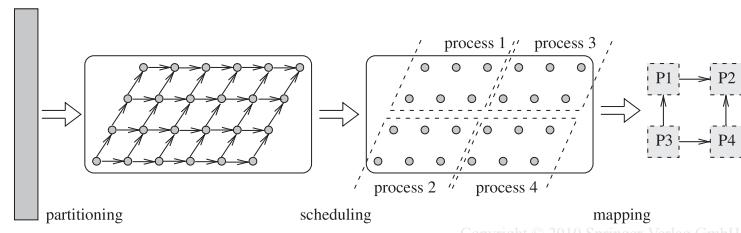
- Decomposition of the algorithm into tasks.
- Specification of task dependencies.
- Tasks include an unrestricted set of computations and
 - access to shared variables (on shared memory systems) or
 - they exchange messages by communication operations (on distributed memory systems)
- **Granularity** of a task: Number of computations performed by a task.

Parallelization of programs (II)

2. Assignment of tasks to processes

- Processes execute tasks successively.
- The aim of the assignment of tasks to processes is to execute nearly the same number of computations by each process, such that a good load balance occurs.
- The assignment of tasks to processes is denoted as scheduling.

3. Mapping of processes to physical processors



Topic overview

- Models for parallel systems
- Parallelization of programs
- **Levels of parallelism**
- Data distribution for arrays
- Information exchange
- Parallel Matrix-Vector Multiplication

Levels of parallelism

- Depending on the level considered, tasks of different **granularity** result
 - Instruction-level parallelism
 - Data parallelism
 - Loop parallelism
 - Functional parallelism

Instruction-level parallelism

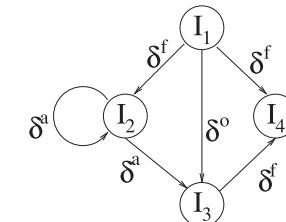
- Multiple instructions can be executed in parallel
- Data dependencies** between instructions I_1 and I_2 limit the parallel execution
 - Flow dependency** (also called *true dependency*): There is a flow dependency from instruction I_1 to I_2 , if I_1 computes a result value in a register or variable which is then used by I_2 as operand.
 - Anti-dependency**: There is an anti-dependency from I_1 to I_2 , if I_1 uses a register or variable as operand which is later used by I_2 to store the result of a computation.
 - Output dependency**: There is an output dependence from I_1 to I_2 , if I_1 and I_2 use the same register or variable to store the result of a computation.

$I_1: R_1 \leftarrow R_2 + R_3$	$I_1: R_1 \leftarrow R_2 + R_3$	$I_1: R_1 \leftarrow R_2 + R_3$
$I_2: R_5 \leftarrow R_1 + R_4$	$I_2: R_2 \leftarrow R_4 + R_5$	$I_2: R_4 \leftarrow R_4 + R_5$
flow dependency	anti dependency	output dependency

Copyright © 2010 Springer-Verlag GmbH

Data dependencies - example

- $I_1: R_1 \leftarrow A$
- $I_2: R_2 \leftarrow R_2 + R_1$
- $I_3: R_1 \leftarrow R_3$
- $I_4: B \leftarrow R_1$



Copyright © 2010 Springer-Verlag GmbH

Data parallelism

- Special constructs to perform the same operation on multiple data elements in parallel (also referred to as SIMD), e.g.,

$$a(1:n) = b(0:n-1) + c(1:n)$$

is the same as

```

for (i = 1:n)
    a(i) = b(i-1) + c(i)
endfor
  
```

Data parallelism

- Data parallelism is often used also in MIMD, e.g., by using the **Single Program Multiple Data** (SPMD) model
 - One parallel program executed in parallel on all processors
 - In practice, many parallel programs are SPMD
 - For example, each processor calculates a part of an array

```

local_size = size/p;
local_lower = me * local_size;
local_upper = (me+1) * local_size - 1;
local_sum = 0;

for (i=local_lower; i<=local_upper; i++)
    local_sum += x[i] * y[i];

Reduce(&local_sum, &global_sum, 0, SUM);
  
```

Loop parallelism

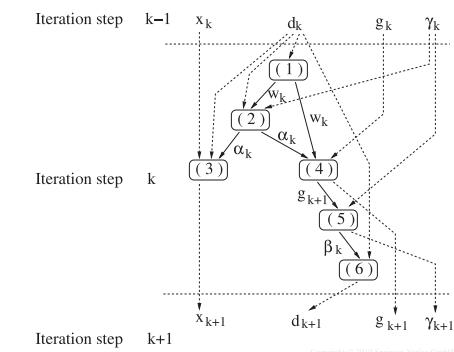
- Execute loop iterations in parallel
 - Example: forall and doall
 - Careful, since they may behave differently (check language etc.)

```
for (i=1:4)           forall (i=1:4)           dopar (i=1:4)
  a(i)=a(i)+1          a(i)=a(i)+1          a(i)=a(i)+1
  b(i)=a(i-1)+a(i+1)  b(i)=a(i-1)+a(i+1)  b(i)=a(i-1)+a(i+1)
endfor                 endforall             enddopar
```

start values	after	after	after
	for-loop	forall-loop	dopar-loop
a(0) 1			
a(1) 2	b(1) 4	5	4
a(2) 3	b(2) 7	8	6
a(3) 4	b(3) 9	10	8
a(4) 5	b(4) 11	11	10
a(5) 6			

Functional parallelism

- Functional parallelism (a.k.a. task parallelism) relies on dividing the program into independent tasks
- The tasks and their dependencies can be represented in a task graph
- Tasks can be executed in parallel as long as their task dependencies are maintained
- Tasks are scheduled either static or dynamically on multiple processors
 - Task pool is a popular model for dynamic task scheduling



Arrangement of tasks / threads

Parallel design patterns (= structures of coordination of the threads) can be used for organizing the cooperation of tasks / threads of a program:

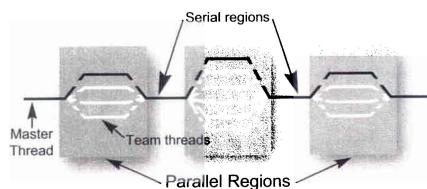
- Creation of threads
- Fork-Join
- Parbegin-Parend
- SPMD and SIMD
- Master-Slave or Master-Worker
- Client-Server-Model
- Pipelining
- Taskpools
- Producer-Consumer-Threads

Creation of threads

- The creation of processes or threads can be carried out *statically* or *dynamically*.
 - Static thread creation:** A *fixed number* of processes or threads are created at program start, all processes or threads exist during the entire execution of the parallel program, and are terminated when program execution is finished.
 - Dynamic thread creation:** Allow creation and termination of processes or threads dynamically at arbitrary points during program execution, i.e., at *run-time*.

Fork-Join

- An existing thread T1 creates a second thread T2, or a group of threads
- Arbitrary nesting
- Specific characteristics in different parallel programming languages and environments



Parbegin-Parend

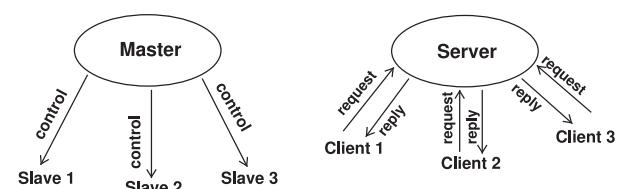
- Simultaneous creation and destruction of several threads (structured variant of thread creation)
- The statements included by Parbegin-Parend are mapped to separate threads;
The statements following the Parend statement are executed after all additional threads are destroyed
- Actual parallel execution depends on implementation
- Specific characteristics in individual programming languages and environments (e.g., parallel sections in OpenMP)

SPMD and SIMD

- SIMD - **S**ingle **I**nstruction, **M**ultiple **D**ata
- SPMD - **S**ingle **P**rogram, **M**ultiple **D**ata
- All threads execute the same program but with different data
 - SIMD: *synchronously*, i.e., all threads execute the same instruction simultaneously. This means data-parallelism in the strict sense.
 - SPMD: *asynchronously*, i.e., at a time different threads execute different program statements at the same time.

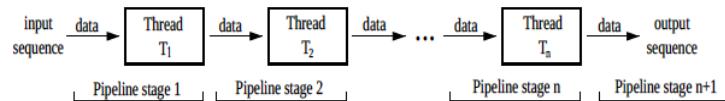
Master-Slave or Master-Worker, vs. Client-Server model

- **Master-Slave:** One single thread controls the whole computations of a program; creates mostly similar Worker- or Slave-threads which get computations assigned
- **Client-Server-Model:** Several client threads make requests to the server thread; server thread handles client requests concurrently/parallel (Extensions: Several server threads or threads which are client and server at the same time)



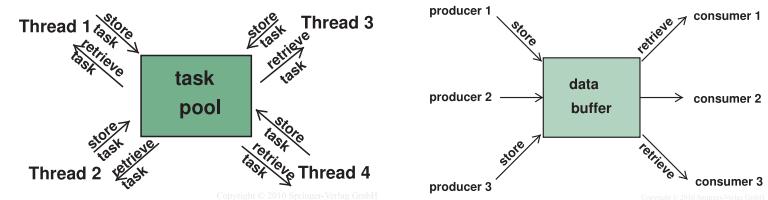
Pipelining

- Threads T_1, \dots, T_n are logically ordered in a specified order;
- Thread T_i gets the output of thread T_{i-1} as input and computes its output, that will be used by thread T_{i+1} , $i = 2, \dots, n - 1$, as input;
- Thread T_1 gets its input from other program parts; T_n provides its output to other program parts
- Parallelism despite of data dependencies



Taskpools and Producer-Consumer

- Taskpool:** Data structure managing program parts in form of *functions* (tasks) that have to be executed; execution by a fixed number of threads, that access the taskpool for extraction and storage of tasks.
- Producer-Consumer:** Producer threads create data and consumer threads use the data; common *data structure* of fixed size of the storage of data.



Topic overview

- Models for parallel systems
- Parallelization of programs
- Levels of parallelism
- Data distribution for arrays**
- Information exchange
- Parallel Matrix-Vector Multiplication

Data distribution for arrays

- Data distribution, data decomposition, and data partitioning: data are partitioned into smaller pieces that are distributed to the processor for execution
 - Distributed memory:** data assigned to a processor are stored in the local memory and can only be accessed by this processor (owner)
 - Shared memory:** data are stored in the same shared memory and processors access different data according to the data distribution pattern → no access conflicts
- In the following:
 - data distribution for one-dimensional arrays
 - data distribution for two-dimensional arrays

Data distribution for one-dimensional arrays

Blockwise data distribution for p processors:

- Array $v = (v_1, \dots, v_n)$ of length n
- Decomposition of array v into p blocks with $\lceil n/p \rceil$ consecutive elements each:
 - Block j contains the consecutive elements with indices: $(j-1) \cdot \lceil n/p \rceil + 1, \dots, j \cdot \lceil n/p \rceil$ for $1 \leq j < p$
 - Block p contains the elements with indices: $(j-1) \cdot \lceil n/p \rceil + 1, \dots, n$
- Block j is assigned to processor j , $(1 \leq j \leq p)$.

Example: For $n = 14$ and $p = 4$

- P1: v_1, v_2, v_3, v_4 ,
- P2: v_5, v_6, v_7, v_8 ,
- P3: $v_9, v_{10}, v_{11}, v_{12}$,
- P4: v_{13}, v_{14} .

blockwise

1	2	3	4	5	6	7	8
P1	P2	P3	P4				

Cyclic data distribution

- Elements of an array $v = (v_1, \dots, v_n)$ are assigned to p processors in a **round robin** way, i.e., v_i is assigned to processor $P_{(i-1) \bmod p + 1}$, $i = 1, \dots, n$
- Processor P_j owns the array elements $j, j + p, \dots, j + p \cdot (\lceil n/p \rceil - 1)$, for $j \leq n \bmod p$
- Processor P_j owns the array elements $j, j + p, \dots, j + p \cdot (\lceil n/p \rceil - 2)$, for $n \bmod p < j \leq p$.

cyclic

1	2	3	4	5	6	7	8
P1	P2	P3	P4	P1	P2	P3	P4

Example: For $n = 14$ and $p = 4$ the following results:

- $n \bmod p = 14 \bmod 4 = 2$
- For $1 \leq j \leq 2$: P_j owns array elements $j, j+4, j+4*2, j+4*(4-1)$
For $2 < j \leq 4$: P_j owns array elements $j, j+4, j+4*(4-2)$
 - P1: v_1, v_5, v_9, v_{13} ,
 - P2: v_2, v_6, v_{10}, v_{14} ,
 - P3: v_3, v_7, v_{11} ,
 - P4: v_4, v_8, v_{12} .

Block-cyclic data distribution

- Combination of the blockwise and the cyclic distribution
- Subdivision of array $v = (v_1, \dots, v_n)$ into blocks of size b ; usually $b \ll n/p$.
- Cyclic distribution of the blocks to the processors P_1, \dots, P_p .

block-cyclic

1	2	3	4	5	6	7	8	9	10	11	12
P1	P2	P3	P4	P1	P2						

Data distribution for two-dimensional arrays

Distribution in only one of the two dimensions

- **blockwise columnwise** (or blockwise rowwise):
 - A block of contiguous columns (or rows) of equal size; block i is assigned to P_i , $i = 1, \dots, p$
- **cyclic columnwise** (or cyclic rowwise):
 - round robin distribution of columns (or rows) to processors

blockwise

1	2	3	4	5	6	7	8
2	P1	P2	P3	P4			
3							
4							

cyclic

1	2	3	4	5	6	7	8
2	P1	P2	P3	P4	P1	P2	P3
3							
4							

Data distribution for two-dimensional arrays

- **block-cyclic columnwise** (or rowwise):

- blocks of contiguous columns (or rows) are assigned to the processors in a cyclic way to processors P_1, \dots, P_p

block-cyclic

	1	2	3	4	5	6	7	8	9	10	11	12
1												
2	P ₁	P ₂	P ₃	P ₄	P ₁	P ₂						
3												
4												

Checkerboard distribution of two-dimensional arrays

- The processors are arranged in a virtual mesh of size $p_1 \times p_2 = p$
- Distribution of the data along both dimensions
- Distribution of the elements of an array of size $n_1 \times n_2$ in blockwise, cyclic, and block-cyclic checkerboard pattern

Checkerboard distribution of two-dimensional arrays

Blockwise checkerboard distribution:

- Decomposition of the array into $p_1 \times p_2 = p$ blocks
- Block (i,j) , $1 \leq i \leq p_1$, $1 \leq j \leq p_2$ contains the elements (k,l) with $k = (i-1) \cdot \lceil n_1/p_1 \rceil + 1, \dots, i \cdot \lceil n_1/p_1 \rceil$ and $l = (j-1) \cdot \lceil n_2/p_2 \rceil + 1, \dots, j \cdot \lceil n_2/p_2 \rceil$.
- Block (i,j) is assigned to processor (i,j) in the processor grid

	1	2	3	4	5	6	7	8
1								
2	P ₁			P ₂				
3								
4	P ₃			P ₄				

Checkerboard distribution of two-dimensional arrays

Cyclic checkerboard distribution:

- Array element (k,l) is assigned to the processor with mesh position $((k-1) \bmod p_1 + 1, (l-1) \bmod p_2 + 1)$.
- The processor at position (i,j) owns all array elements (k,l) with $k = i + s \cdot p_1$ and $l = j + t \cdot p_2$ for $0 \leq s < n_1/p_1$ and $0 \leq t < n_2/p_2$.

	1	2	3	4	5	6	7	8
1	P ₁	P ₂						
2	P ₃	P ₄						
3	P ₁	P ₂						
4	P ₃	P ₄						

Checkerboard distribution of two-dimensional arrays

Block-cyclic checkerboard distribution:

- Decomposition of the array into blocks of size $b_1 \times b_2$
- Array element (m,n) belongs to block (k,l) with $k = \lceil m/b_1 \rceil$ and $l = \lceil n/b_2 \rceil$.
- Block (k,l) is assigned to processor at mesh position $((k-1) \bmod p_1 + 1, (l-1) \bmod p_2 + 1)$.
- Special cases:
 $b_1 = b_2 = 1$ cyclic checkerboard distribution
 $b_1 = n_1/p_1$ and $b_2 = n_2/p_2$ blockwise checkerboard distribution

1	2	3	4	5	6	7	8	9	10	11	12
1	P ₁	P ₂	P ₁	P ₂	P ₁	P ₂					
2											
3	P ₃	P ₄	P ₃	P ₄	P ₃	P ₄					
4											

Topic overview

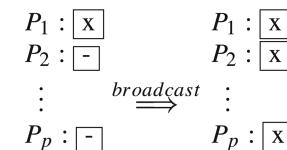
- Models for parallel systems
- Parallelization of programs
- Levels of parallelism
- Data distribution for arrays
- Information exchange
- Parallel Matrix-Vector Multiplication

Information exchange

- The **information exchange** between processors of a parallel systems depends on the **organization of the memory subsystem**:
 - shared address space: shared variables
 - distributed address space: explicit communication operations
- **Shared variables:** Concurrent accesses through several processor to the same address is protected by *synchronization operations*
 - Serialization of concurrent accesses
 - Prevention of race conditions
 - Simple synchronization by using (lock/unlock).
- **Communication operations:** Information exchange by **sending messages** (message passing);
 - Differentiation between *point-to-point* communication and *global* communication

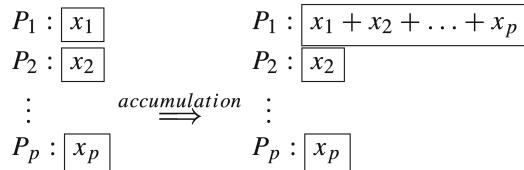
Overview of communication operations: Send-Receive / Broadcast

- **Point-to-Point transfer:** A processor P_i (sender) sends a message to another processor P_j (receiver).
 - The **sender** executes a send operation (with the specification of a *send buffer*) and with the identification number of the receiver.
 - The **receiver** executes a corresponding receive operation with the specification of an *receive buffer* and with the specification of the identification number of the sender.
- **Single Broadcast:** A specific processor P_i (root) sends the same message to all other processors. Depiction:



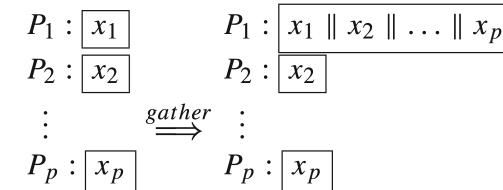
Overview of communication operations: Reduction / single-accumulation

- Single-Accumulation Operation:** Each process sends a message to a specific processor P_i (root) with data of the same type.
- The messages are combined elementwise with a specific **reduction operation**.
→ the result on the root process P_i is a single (composed) message.
- Each process specifies a buffer with the data to be combined and the reduction operation which to be used.



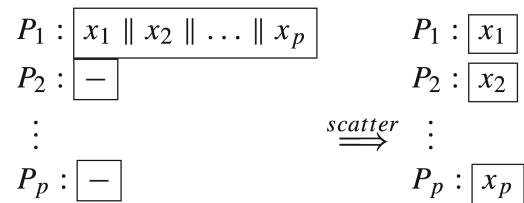
Overview of communication operations: Gather

- Gather:** Each process sends to a specific processor (root) a message. The root processor collects the messages *without any reduction*.
- Each process specifies a buffer storing the data to be sent. The root process specifies an *additional* buffer for the collected messages.



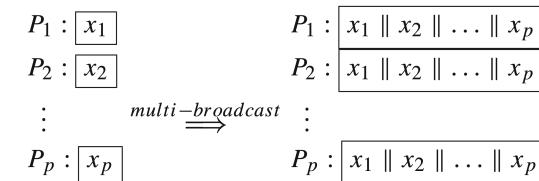
Overview of communication operations: Scatter

- Scatter:** A specific processor P_i (root) sends to the other processors a message, which might be different for each receiver.



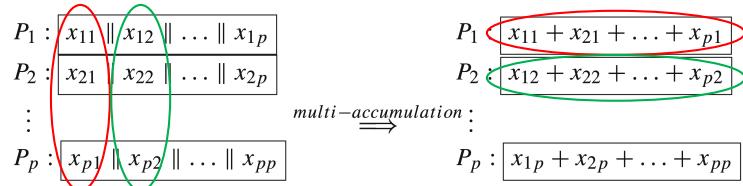
Overview of communication operations: Multi-Broadcast

- Multi-Broadcast Operation:** Each processor executes a single broadcast operation, i.e., *each other* processor the *same message*.
 - Contrary, each processor receives a message from each other processor, where the different receivers receive from one sender the same message.
 - Note: There is *no specific* root process



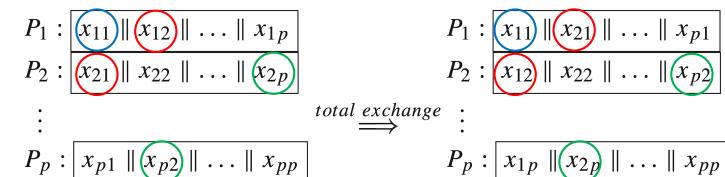
Overview of communication operations: Multi-Accumulation

- Multi-Accumulation Operation:** Each processor executes a single accumulation operation, i.e., each process provides for each other process a possible different message.
 - The messages specific for each receiver are combined with a *reduction operation*, such that each receiver gets a combined message.
 - Note: There is *no specific root process*



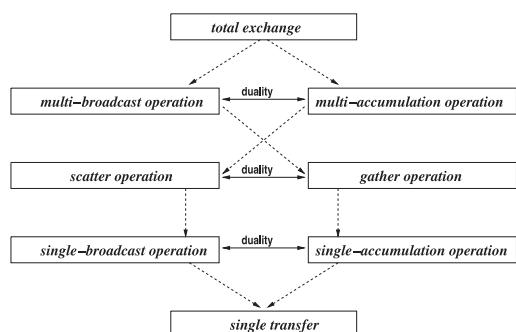
Overview of communication operations: Total Exchange

- Total Exchange:** Each processor sends to each other processor a possibly *different message*, without using a reduction operation, i.e., each processor executes a scatter operation.
 - Contrary, each processor receives from each other processor a possibly different messages, i.e., each processor executes a gather operation.
 - Note: There is *no specific root process*



Hierarchy of Communication Operations

- The communication operations result from a *stepwise specialization* from the most general operation (total exchange)
 - Representation as a hierarchy is possible:



Topic overview

- Models for parallel systems
- Parallelization of programs
- Levels of parallelism
- Data distribution for arrays
- Information exchange
- Parallel Matrix-Vector Multiplication

Parallel Matrix-Vector Multiplication

► Multiplication of

- a dense $n \times m$ -matrix $\mathbf{A} \in \mathbb{R}^{n \times m}$, $\mathbf{A} = (a_{ij})_{i=1, \dots, n, j=1, \dots, m}$
- and a vector $\mathbf{b} \in \mathbb{R}^m$, $\mathbf{b} = (b_1, \dots, b_m)$
- with result vector $\mathbf{c} = (c_1, \dots, c_n) \in \mathbb{R}^n$

$$c_i = \sum_{j=1}^m a_{ij} b_j, \quad i = 1, \dots, n,$$

► There exist two implementation variants differing in the loop order over i and j , $i, j = 1, \dots, n$.

- computation of n scalar products.
- linear combination of columns

H. Grahn/DV2597/DV2606

Matrix-Vector Product using scalar products

- Computation of n scalar products (vector-vector-multiplication of rows $\mathbf{a}_1, \dots, \mathbf{a}_n$ von \mathbf{A} with vector \mathbf{b}):

$$\mathbf{A} \cdot \mathbf{b} = \begin{pmatrix} (\mathbf{a}_1, \mathbf{b}) \\ \vdots \\ (\mathbf{a}_n, \mathbf{b}) \end{pmatrix},$$

- A scalar product is defined as:
 $(\mathbf{x}, \mathbf{y}) = \sum_{j=1}^m x_j y_j$
 for two vectors $\mathbf{x}, \mathbf{y} \in \mathbb{R}^m$ with $\mathbf{x} = (x_1, \dots, x_m)$ and $\mathbf{y} = (y_1, \dots, y_m)$

- Corresponding sequential algorithm in C-notation:
`for (i=0; i<n; i++) c[i] = 0;
 for (i=0; i<n; i++)
 for (j=0; j<m; j++)
 c[i] = c[i] + A[i][j] * b[j];`
 with a two-dimensional array A and one-dimensional array b, c . (The indices start with 0 as usual in C).

Matrix-Vector Product based on linear combinations

- Computation of a linear combination of columns $\tilde{\mathbf{a}}_1, \dots, \tilde{\mathbf{a}}_m$ of \mathbf{A} , with coefficients (b_1, \dots, b_m) , i.e.

$$\mathbf{A} \cdot \mathbf{b} = \sum_{j=1}^m b_j \tilde{\mathbf{a}}_j.$$

- Corresponding sequential algorithm in C-notation:
`for (i=0; i<n; i++) c[i] = 0;
 for (j=0; j<m; j++)
 for (i=0; i<n; i++)
 c[i] = c[i] + A[i][j] * b[j];`
 ► For each $j = 0, \dots, n-1$ a column $\tilde{\mathbf{a}}_j$ is added to the linear combination.
 ► This sequential program is equivalent to the previous one, since the loops over i and j can be exchanged due to data independence.

Parallel Matrix vector-Product

The two sequential representations give rise to two different parallel implementations

- (a) Row-oriented representation of matrix A and the computation of n scalar products:
 Parallel implementation in which each processor computes about n/p scalar products (p denotes the number of processors)
- (b) Column-oriented representation of matrix A and computation of a linear combination:
 Parallel implementation in which each processor computes a part of the linear combination using about n/p column vector

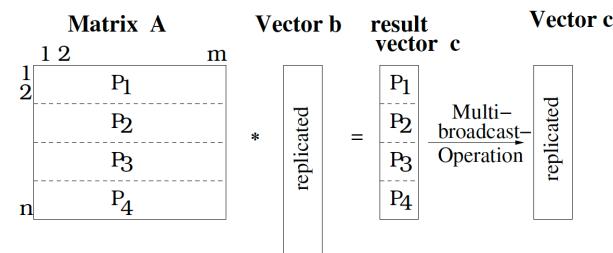
H. Grahn/DV2597/DV2606

Matrix-Vector Multiplication: distribution of rows (1)

- ▶ Matrix-vector multiplication based on **scalar products**
- ▶ Set of p processors P_k , $k = 1, \dots, p$ with **distributed memory**
- ▶ Each processor computes that scalar product for which it own the corresponding row of matrix **A**.
- ▶ Data distribution:
 - ▶ Row-oriented blockwise distribution of matrix **A**:
 - Processor P_k stores the rows a_i for $i = n/p \cdot (k-1) + 1, \dots, n/p \cdot k$ in its local memory, $k = 1, \dots, p$
 - ▶ Result vector $c = (c_1, \dots, c_n)$ has a block wise distribution.

H. Grahn/DV2597/DV2606

Matrix-Vector-Multiplication: distribution of row (2)



- ▶ When the matrix-vector-product is used within a large algorithm like iteration methods, a certain distribution of c might be required. Example: Vector c should have the same distribution as b .
 - ▶ Each processor P_k , $k = 1, \dots, p$ sends its block $(c_{n/p \cdot (k-1)+1}, \dots, c_{n/p \cdot k})$ to all other processors by a **multi-broadcast** operation.

Matrix-Vector Multiplication: distribution of rows (3)

- ▶ Parallel implementation SPMD program for processor P_k mit $k = 1, \dots, p$
- ▶ Row-wise distribution
 - ▶ Each processor has a local array `local_A` of size `local_n × m`
 - ▶ Processor P_k stores the following data:

$$\text{local_A}[i][j] = A[i + (k-1) * n/p][j]$$

- ▶ Result vector: local array `local_c` of size `local_n`
- ▶ After the multi-broadcast operation

$$c[i + (k-1) * n/p] = \text{local_c}[i]$$

H. Grahn/DV2597/DV2606

Matrix-Vector-Multiplication: distribution of rows (4)

- ▶ Program fragment in C-notation and MPI communication operation

```
local_n = n/p;
for (i=0; i<local_n; i++) local_c[i] = 0;
for (i=0; i<local_n; i++)
  for (j=0; j<m; j++)
    local_c[i] = local_c[i] + local_A[i][j] * b[j];
MPI_Allgather(local_c,local_n,MPI_DOUBLE,
              global_c,local_n,MPI_DOUBLE,comm);
```

H. Grahn/DV2597/DV2606

Parallel Matrix-Vector-Multiplication: Shared Memory

- ▶ SPMD program with a distribution of the computations
- ▶ No explicit distribution of data
- ▶ Each process accesses a different part of the matrix A
- ▶ Each process computes n/p components of the result vector c and uses the corresponding n/p rows of matrix A
→ No access conflict

Parallel Matrix-Vector-Multiplication: Shared Memory

- ▶ SPMD program with private variable k holding the processor ID.

```
local_n = n/p;
for (i=0; i<local_n; i++) c[i+(k-1)*local_n] = 0;
for (i=0; i<local_n; i++)
    for (j=0; j<m; j++)
        c[i+(k-1)*local_n] =
            c[i+(k-1)*local_n] + A[i+(k-1)*local_n][j] * b[j];
synch();
```